

# Generative Adversarial Networks (GANs)

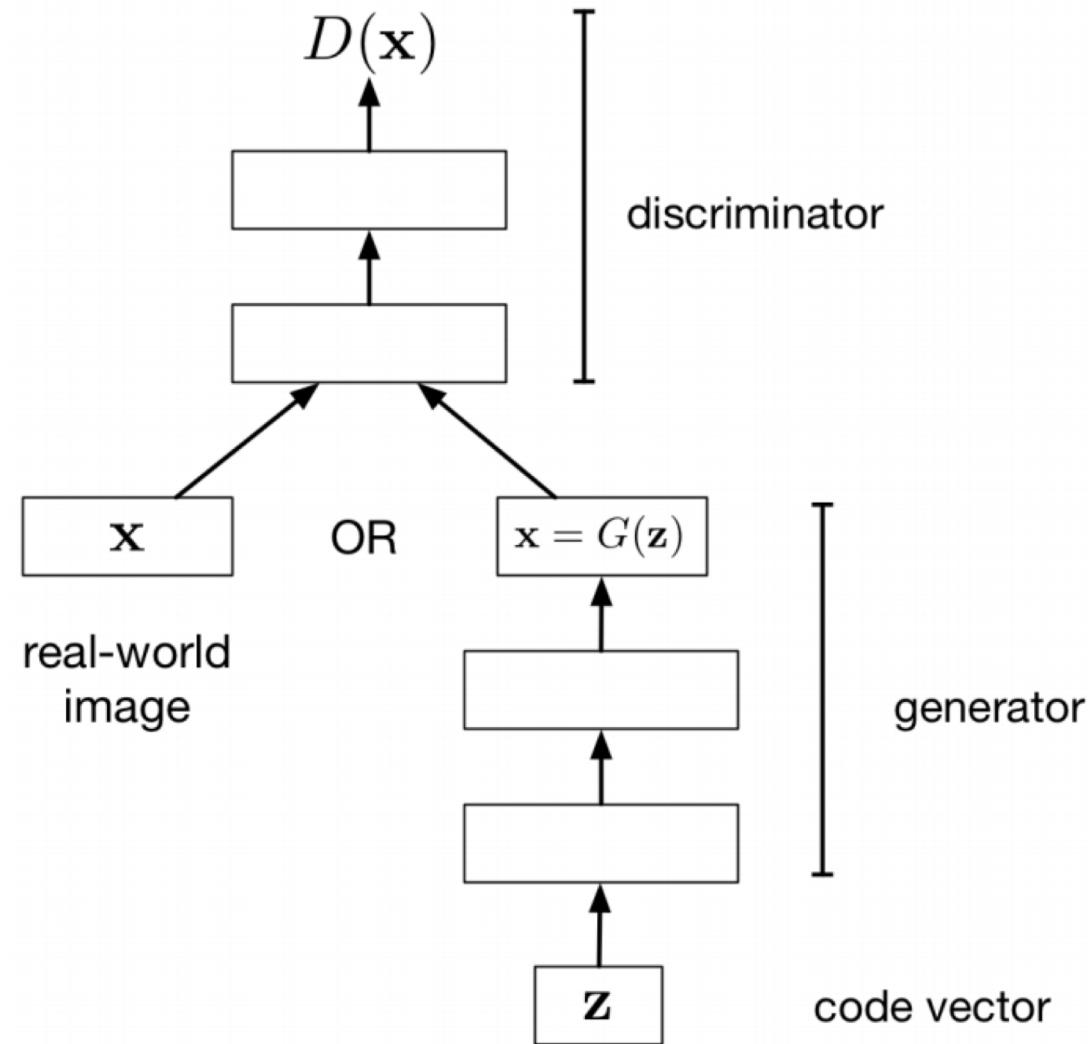
# GANs

- Generative
  - Generates samples instead of discriminates
    - Models  $p(x)$  instead of  $p(y|x)$
- Adversarial
  - Trained in an “Adversarial” setting (more on this to come)
- Networks
  - Using a neural network

# GANs Overview

- Implicit data generation – create samples without any explicit data
- If we have some mechanism of evaluating the quality of implicitly generated data, then we can compute a gradient for training a network
- The big idea of GANs is we train two networks simultaneously –
  - The generator network generates real looking samples
  - The discriminator network tries to figure out whether or not a sample from came from the generator network or a training set
- The generator network tries to fool the discriminator

# Adversarial Training



# GAN Training

- Discriminator – Predicts the probability of being true data
- Generator – Creates samples from the distribution of interest
- Two cost functions –
  - Discriminator cost function:
$$J_D = \mathbb{E}_{x \sim \mathcal{D}}[-\log(D(x))] + \mathbb{E}_z \left[ -\log(1 - D(G(z))) \right]$$
  - Generator cost function:
$$J_G = -J_D$$

# Minimax Formulation of GAN Training

- Discriminator cost function:

$$J_D = \mathbb{E}_{x \sim \mathcal{D}}[-\log(D(x))] + \mathbb{E}_z \left[ -\log(1 - D(G(z))) \right]$$

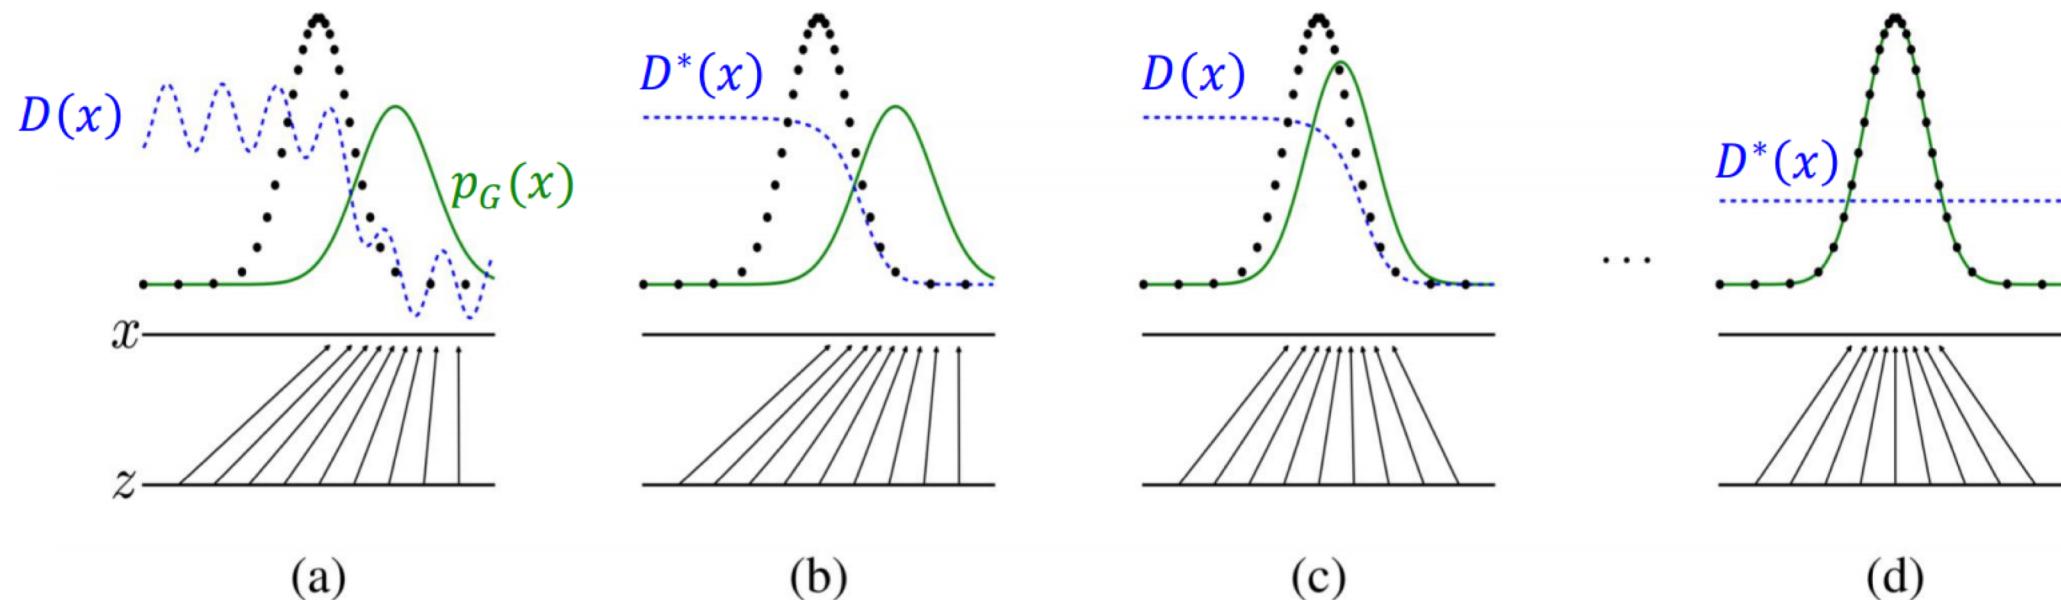
- Generator cost function:

$$J_G = -J_D$$

- Minimax Formulation:

$$\max_G \min_D J_D$$

# Training the GAN



# Problem with the Minimax Formulation

- Since  $J_G = -J_D$

$$J_G = \mathbb{E}_{x \sim \mathcal{D}} [\log(D(x))] + \mathbb{E}_z [\log(1 - D(G(z)))]$$

- There is a problem when the discriminator is much better than the generator

$$\begin{aligned} D(G(z)) &= 0 \\ \mathbb{E}_z [\log(1 - D(G(z)))] &= 0 \end{aligned}$$

- There is no gradient to improve the generator

# A Better Cost Function

- Original cost function –

$$\begin{aligned} J_G \\ = \mathbb{E}_{x \sim \mathcal{D}} [\log(D(x))] \\ + \mathbb{E}_z [\log(1 - D(G(z)))] \end{aligned}$$

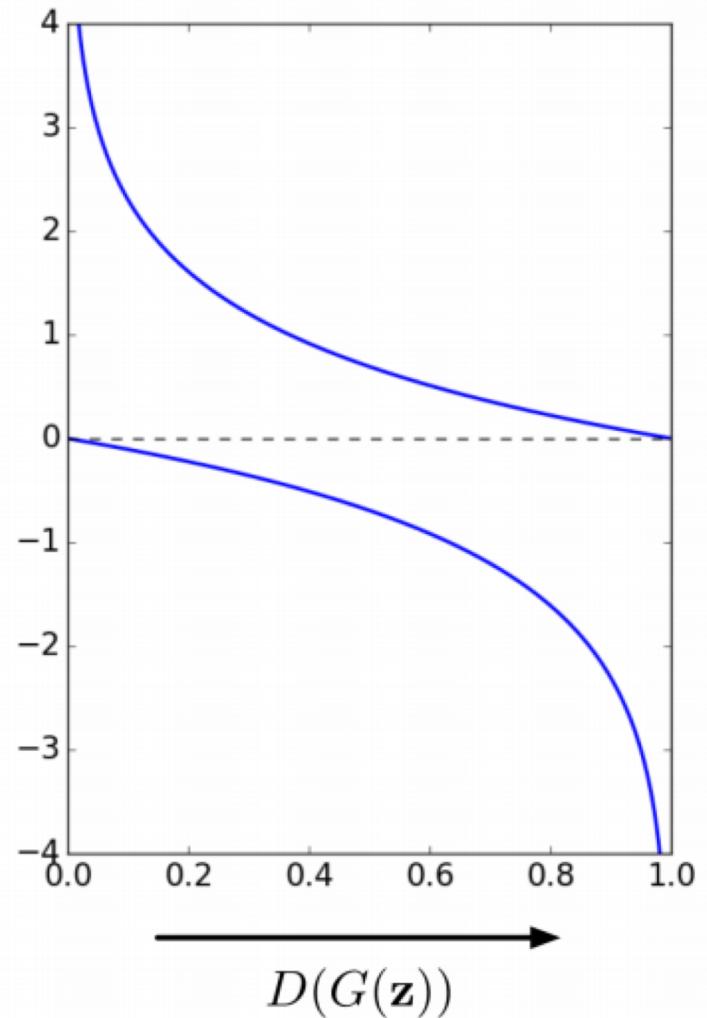
- New cost function –

$$J_G = \mathbb{E}_z [\log(D(G(z)))]$$

- Change is to the target class

modified  
cost

minimax  
cost



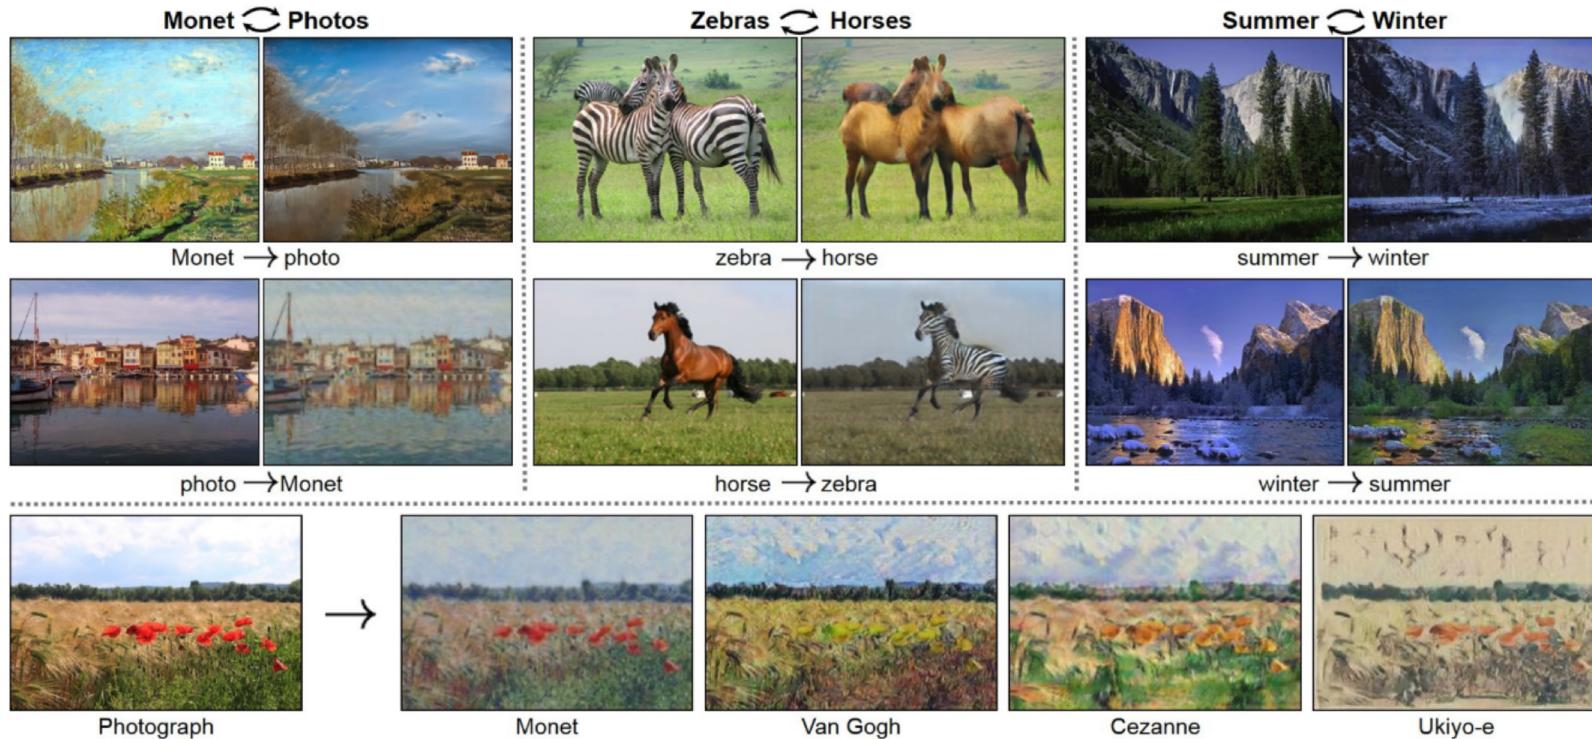
(how well it fooled  
the discriminator)

# GANs

- Since their original formulation in 2014, hundreds of papers have introduced various training techniques and architectures
- Many common architectures focus on image generation and rely on the DC-GAN architecture
  - De-Convolution Generative Adversarial Network
  - Generator and Discriminator are convolutional neural networks

# CycleGAN

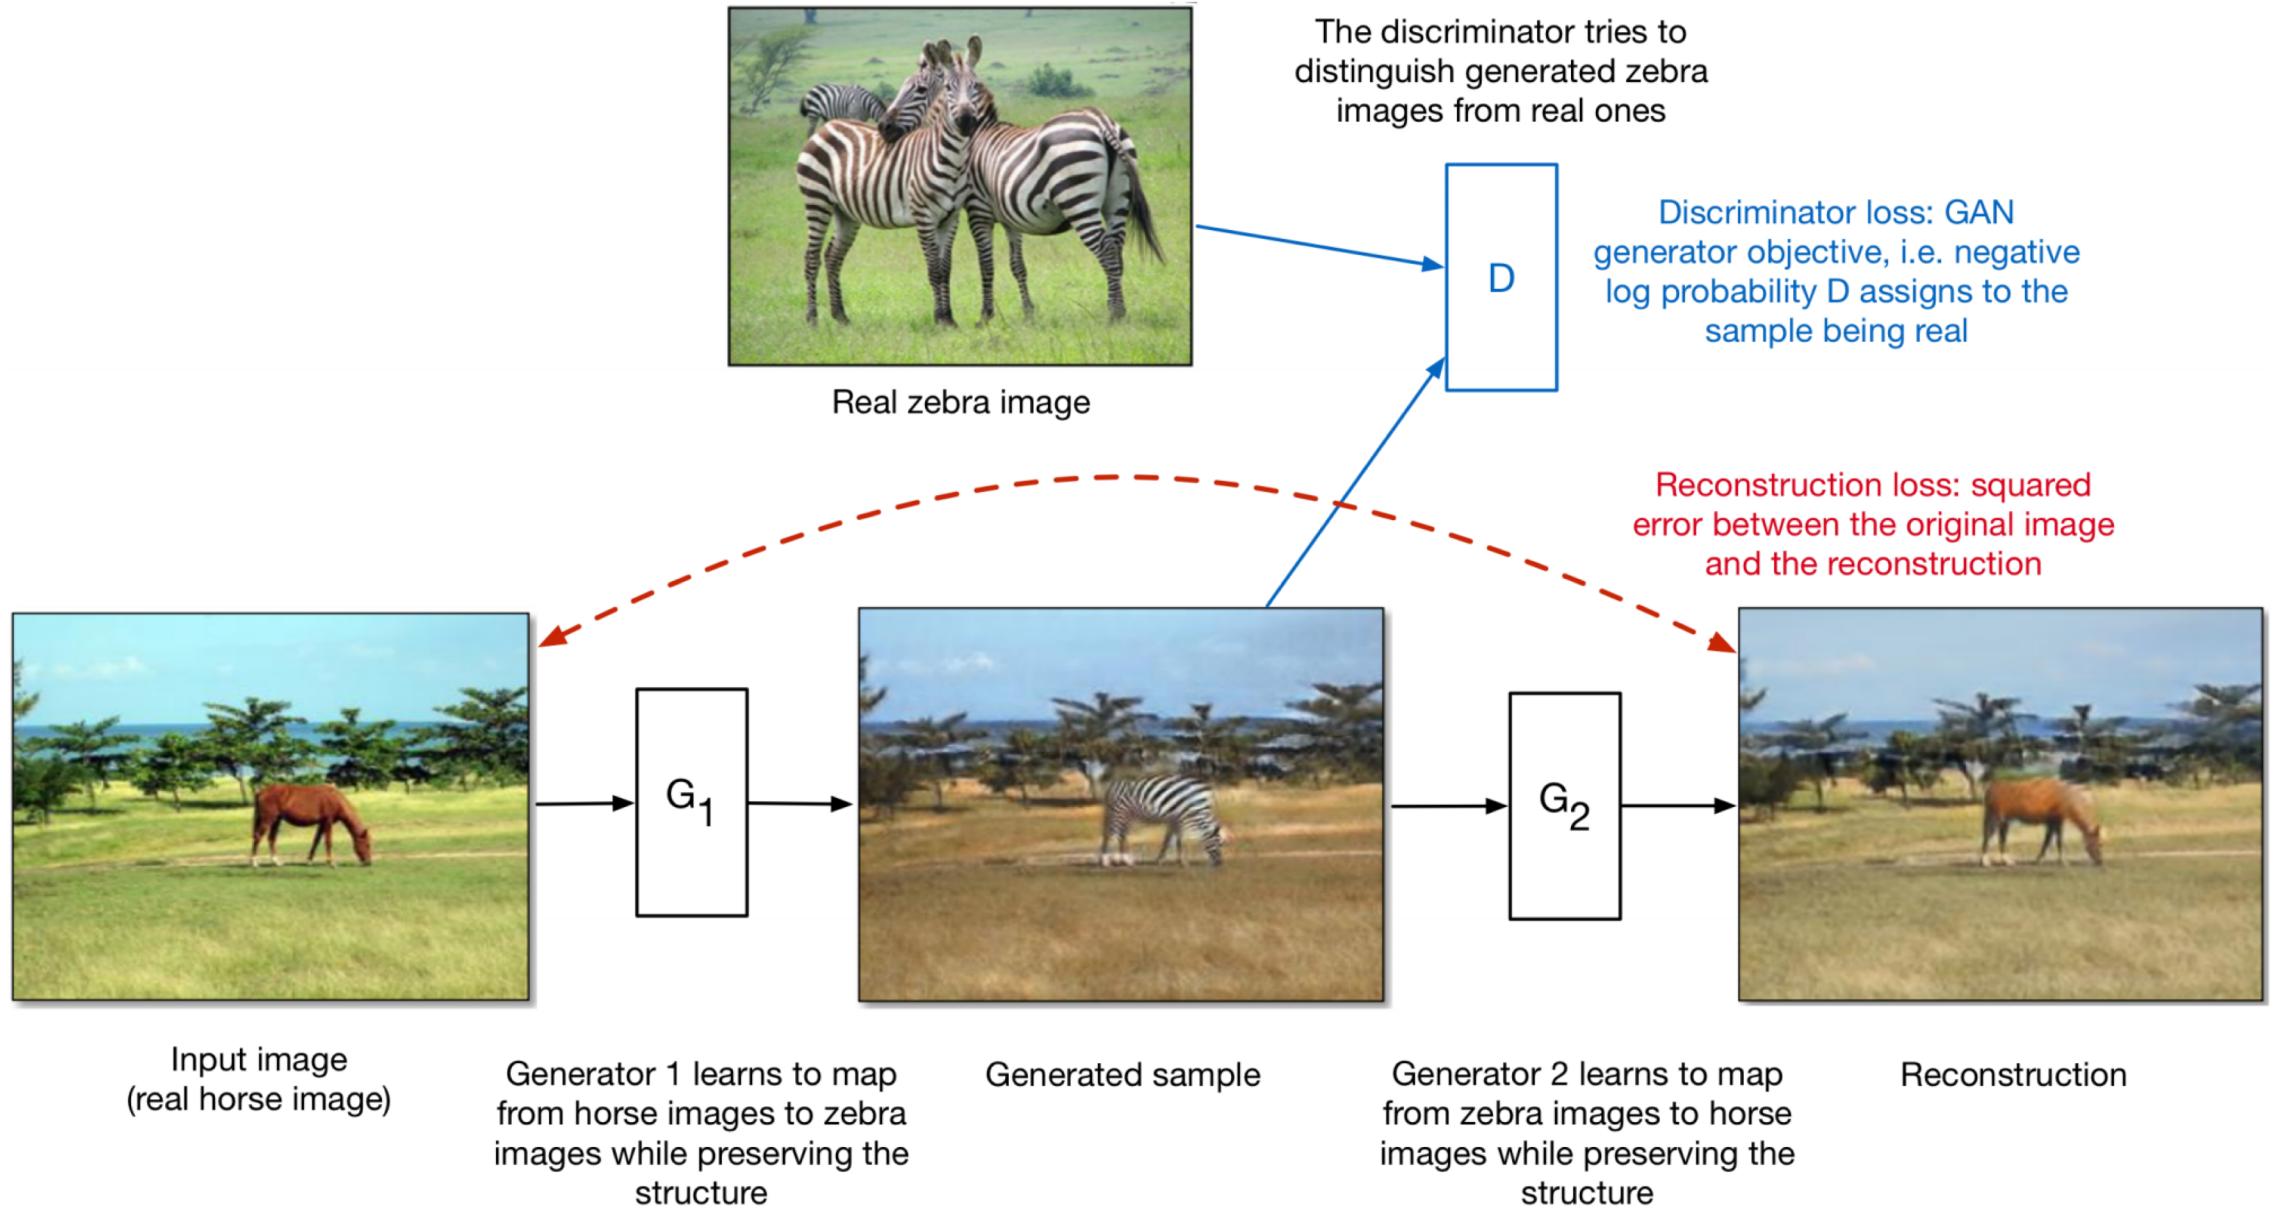
Style Transfer: Change the style of an image while preserving its content



Data: Two **unrelated** collections of images, one set for each style

# CycleGANs

- If we have paired data, this just becomes a supervised learning problem
  - Often times, it is hard or impossible to find paired data
- The CycleGAN architecture learns style transfer from unpaired data
  - Train two different generator networks to go from style 1 to style 2 and vice versa
  - Make sure the generated samples of style 2 are indistinguishable from real images in the generator set
  - Make sure the generators are cycle consistent – mapping from style 1 to style 2 and back should give you back almost the original image



Total loss = discriminator loss + reconstruction loss

# CycleGAN – Style Transfer Between Photos and Maps

